

Handout: A Quick Summary on Undecidability and NP-Completeness

CS6033 Design and Analysis of Algorithms I

By Prof. Yi-Jen Chiang

CSE Dept., Tandon School of Engineering, New York University

Note: There are 5 pages.

Complementary to designing efficient algorithms, you should know that there are problems that cannot even be computed/solved/decided by a computer (the topic of **undecidability**), and that there are problems that can be solved by a computer but **very unlikely to be solved/computed efficiently** (the topic of **NP-completeness**). This handout gives a quick summary on these topics below.

1. Undecidability

Turing Machine

Turing machine is a mathematical model of computation. It consists of a finite-state control unit, a tape, and a read/write head pointing at some tape cell from the control unit (see Fig. 1). The control unit operates in discrete steps; at each step it reads the symbol from the read/write head, and depending on the current state q and the symbol x read, it does the following:

1. Put the control unit into a new state q' .
2. *Either*
 - (a) Write a symbol y in the tape cell of the read/write head position, effectively replacing x by y ; *or*
 - (b) Move the read/write head one tape cell to the left or right.

The action can be expressed as a *rule of action* $(q, x) \rightarrow (q', y)$ or $(q, x) \rightarrow (q', m)$ where m is L (denoting moving left) or R (denoting moving right).




Figure 1: An example of a Turing machine.

Initially, the control unit is in the *initial state* s , and the input string w is given on the tape with the read/write head pointing to the leftmost symbol of w . Then the Turing machine acts according to the rules of action; when it reaches a special state, the *halting state* h , the computation stops and

the result of the computation is left on the tape. It is possible that for some input the computation never stops.

Finite-Length Encoding of a Turing Machine. Note that there are **fixed** number of states and also **fixed** number of symbols that can appear on the tape. Therefore, the number of rules of action is also **fixed** (e.g., if there are 6 states and 3 symbols then there are 18 rules, since the number of tuples (q, x) where q is a state and x is a symbol is $6 \cdot 3$). In other words, each Turing machine has a **finite description**, and thus can be **encoded into a finite-length string**.

Turing Machines and Functions. We can see that each Turing machine TM corresponds to a function F : for any input string w , if TM stops and produces the output α , then w is in the domain of F with $F(w) = \alpha$; for other cases (TM does not stop on w or TM stops but does not accept w as a valid input, etc.) then w is not in the domain of F . It turns out that Turing machines have the same computing powers as computers (in terms of what functions can be computed on them, **not** in terms of the computing efficiencies).

Universal Turing Machine. A nice property is the **universality** of Turing machines, namely, we can have a **universal Turing machine** TM^* : given an input $\langle e(M), w \rangle$ where $e(M)$ is the *encoding* of some Turing machine M and w is a string, TM^* will **simulate** the actions of M on the input w . This essentially means that the universal Turing machine TM^* is a **general-purpose computer** and we can **program** it — the finite-length encoding $e(M)$ is some finite-length **program** (say in C++ or Java, etc.) and the simulation is to run the program on the input w .

Countable vs. Uncountable Sets

We say that a set A is **countable** if **every** element $a \in A$ can be mapped to a **distinct** number from the set \mathbb{N} of natural numbers $\{1, 2, 3, \dots\}$. For example, the set \mathbb{Z} of all integers is countable, since we can count them as follows: count/map $0, 1, -1, 2, -2, \dots$, in that order, as $1, 2, 3, 4, \dots$ (note: counting the integers in the order of $0, 1, 2, 3, \dots$ does *not* work, since this would only go along the positive integers and miss the negative ones). The set \mathbb{Q}^+ of positive rational numbers is countable, since such numbers can be represented as p/q where $p = 1, 2, 3, \dots$ and $q = 1, 2, 3, \dots$ (think of listing p/q in a 2-dimensional table where one dimension is indexed by p and the other by q). We start by counting the tuples (p, q) in the order of $(1, 1), (1, 2), (2, 1), (1, 3), (2, 2), (3, 1), \dots$. That is, with $p + q = 2$ there is one tuple $(1, 1)$; with $p + q = 3$ there are two tuples $(1, 2), (2, 1)$; with $p + q = 4$ there are three tuples $(1, 3), (2, 2), (3, 1)$, etc. (this is called *dovetailing*). In this way all tuples (p, q) are counted and no one is missed. The set \mathbb{Q}^- of all negative rational numbers can be counted symmetrically. Therefore, the set \mathbb{Q} of all rational numbers is countable, since we can first count 0, then the first number in \mathbb{Q}^+ , the first number in \mathbb{Q}^- , the second number in \mathbb{Q}^+ , the second number in \mathbb{Q}^- , and so on. For an infinite set, we say that it is **countably infinite** if it is countable; otherwise it is **uncountably infinite**.

Claim 1. *The set of all Turing machines is countably infinite.*

Proof: Each Turing machine M corresponds to its encoding $e(M)$, which is a finite-length string. Therefore we can sort these encodings/strings first by their length and then by their lexicographical order within the group of the same length. We count these encodings/strings in this final sorted order as $1, 2, 3, \dots$. Clearly all strings are counted and thus the set is countably infinite. \square

Claim 2. *The set of all real numbers in the range $[0, 1)$ is uncountably infinite.*

Proof: The proof uses a proof technique called **diagonalization**. Consider representing such real numbers in binary, then each such $r \in [0, 1)$ is in the form $0.b_1 b_2 b_3 \dots$, where b_i is the i -th bit after

the decimal point, and b_i is either 0 or 1. Assume that the set in question is countable, i.e., we can label all real numbers in $[0, 1)$ as r_1, r_2, r_3, \dots . Then we can list them in the binary form as shown in Fig. 2, where each bit b_{ij} is either 0 or 1. We can assume that each r_i has infinite number of bits,

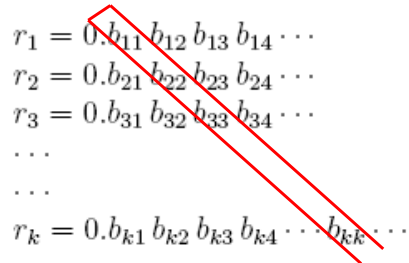
$$\begin{array}{l} r_1 = 0.b_{11} b_{12} b_{13} b_{14} \dots \\ r_2 = 0.b_{21} b_{22} b_{23} b_{24} \dots \\ r_3 = 0.b_{31} b_{32} b_{33} b_{34} \dots \\ \dots \\ \dots \\ r_k = 0.b_{k1} b_{k2} b_{k3} b_{k4} \dots b_{kk} \dots \end{array}$$


Figure 2: Proof by diagonalization.

by appending bits of 0 if needed. Now consider the real number $r^* = 0.b_{11} b_{22} b_{33} b_{44} \dots b_{kk} \dots$, where $b_{ii}, i = 1, 2, \dots$ are the diagonal bits marked by the red lines in Fig. 2. Finally, we let $\overline{r^*} = 0.\overline{b_{11}} \overline{b_{22}} \overline{b_{33}} \overline{b_{44}} \dots \overline{b_{kk}} \dots$, where $\overline{b_{ii}}$ means the complement of b_{ii} : if $b_{ii} = 0$ then $\overline{b_{ii}} = 1$; else ($b_{ii} = 1$) then $\overline{b_{ii}} = 0$. Clearly, $\overline{r^*}$ **differs from each** r_1, r_2, \dots **at the diagonal bit** b_{ii} . However, $\overline{r^*}$ is a real number in $[0, 1)$ and we must have $\overline{r^*} = r_k$ for some $k \in \mathbb{N}$, which is a contradiction since $\overline{r^*}$ differs from r_k at b_{kk} . Therefore it is not possible to label all real numbers in $[0, 1)$ as r_1, r_2, \dots , meaning that the set in question is uncountable. \square

Note that the above proof shows that an uncountable set S has “more elements” than a countable set — when trying to count/label the elements of S by $1, 2, 3, \dots$ there is always some element of S “missing” that cannot be so counted/labeled.

Take-Away Message: There are more functions than Turing machines

For each real number $x \in [0, 1)$ we can define a **distinct function** $f_x : [0, 1) \rightarrow [0, 1)$ as follows: $f_x(0) = x$, and $f_x(t) = t$ for all $t \in [0, 1)$ and $t \neq 0$ (note that each f_x maps the same 0 to a **distinct** value x , and thus each f_x is **distinct**). Since by Claim 2 there are uncountably many real numbers $x \in [0, 1)$, there are also **uncountably many** functions f_x . However, by Claim 1, there are only countably many Turing machines. This means that there are **more functions than Turing machines**, and thus **there are functions that cannot be computed by Turing machines/computers**, i.e., there are problems that cannot be computed/solved/decided by a computer.

A famous undecidable/unsolvable problem is the **halting problem**: “Given a Turing machine encoding/computer program M and an input w , can M terminate on w ?” (For a general M and a general w , this is undecidable! This can be proved by diagonalization.)

2. NP-Completeness

Now we consider the problems that can be computed/solved/decided by Turing machines/computers, but the focus is on whether they can be done efficiently or not. In general, polynomial-time solutions are considered efficient, and solutions requiring more than polynomial time (e.g., exponential time) are considered inefficient.

P vs. NP

Definition of P : The class P is define to be the set of problems that can be solved in polynomial time

(P stands for polynomial time). There is also a class NP , where NP stands for **non-deterministic polynomial time**. For this, we first need to revisit Turing machines. In the Turing machine discussed in the beginning of Sec. 1, each rule of action $((q, x) \rightarrow (q', y)$ or $(q, x) \rightarrow (q', m))$ describes a **deterministic** action, i.e., for each tuple (q, x) there is a unique action. Such Turing machine is called a **deterministic Turing machine**, and for each input w , the computation process is a single path. For a problem in P , the computation is a **single path with polynomial length** for a deterministic Turing machine.

In a **non-deterministic Turing machine**, each rule of action is of the form $(q, x) \rightarrow S$, where S is a **finite set** of possible actions, i.e., $S = \{s_1, s_2, \dots, s_k\}$ for some fixed k and each s_i is either some (q', y) or some (q', m) . This means that for a given state q and the input symbol x read, it will **non-deterministically** choose some $s_i \in S$ as the action (note that we do not say *randomly* since there is **no probability involved**). For each input w , the computation process is a **tree** where each internal node has fan-out $|S|$ (the number of actions in S) if the corresponding rule of action is $(q, x) \rightarrow S$. The computation tree contains many paths; some may not even stop. Among these paths, if **there exists a path that stops in polynomial time** (i.e., **if there exists a path of polynomial length**), then we say that the input w takes **non-deterministic polynomial time**.

Definition of NP : The class NP is defined to be the set of problems that can be solved in non-deterministic polynomial time. Equivalently, it is the set of problems for which **a given solution can be verified in polynomial time**. For example, given a (large) integer N , factoring it is difficult, but given two factors a and b as an answer, we can verify that $N = a \cdot b$ in polynomial time. Intuitively, P is a class of problems that are easy to compute, and NP is a class of problems that are easy to verify the solutions.

It is trivial to see that $P \subseteq NP$: from the deterministic Turing machine whose computation is a single path p of polynomial length, we can construct a non-deterministic Turing machine whose computation tree contains the path p .

All problems in NP have exponential time solutions. Recall that for any problem in NP , the corresponding computation tree of a non-deterministic Turing machine NTM has a path p of polynomial length ℓ . Then we can use a deterministic Turing machine TM to simulate NTM by exploring the computation tree in a **locked-step** fashion: in each round we advance the first path one more step, then the second path one more step, etc.; after the last path is advanced one more step, we repeat the process for the next round. This essentially performs a BFS in the computation tree; when we reach the leaf of path p then we find the solution and can stop. The total time is the **size of the computation tree cut at depth ℓ** ; if the maximum fan-out in the tree is k then such size is $O(k^\ell)$, i.e., the total time is exponential.

Optimization vs. Decision Problems. An optimization problem asks for a solution that *minimizes/maximizes* an objective function, while a decision problem asks for a yes/no question. Note that the **optimization** version of a problem is **at least as hard as** the corresponding decision version. For example, for the shortest-path problem, the optimization version asks for the length of a shortest path from s to t , while the decision version asks, for a given value k , whether there is a path from s to t with length at most k . Clearly, if we have the optimal length ℓ , we can answer yes/no by checking whether $\ell \leq k$.

To show that a problem is hard, it suffices to show that its **decision version is already hard** (then the optimization version is even harder). In the following, we only consider the decision problems.

NP-Complete Problems: The **hardest** problems in NP

Definition of NP-Completeness: We say that a problem Q is **NP-complete** if

(a) Q is in NP , and

(b) for any problem K in NP , K can be *reduced* (i.e., *transformed*) to Q in polynomial time.

Meaning of (b): If there is a polynomial-time solution A to Q , then all problems K in NP can be solved in polynomial time —

method:

reduce/transform K to Q (which takes polynomial time), then solve Q in polynomial time by A , which then tells the answer to K (if yes to Q then yes to K , and if no to Q then no to K).

Therefore K can be solved in polynomial time.

Take-Away Message: Status of P vs. NP

So far, we believe that $P \neq NP$, and we try to use NP-complete problems as witnesses that $P \neq NP$, as follows.

There are thousands of NP-complete problems; none of them has a polynomial-time solution (only exponential-time solutions are known), and if any of them has a polynomial-time solution, then all of them would have polynomial-time solutions (this follows from **Meaning of (b)** above), and thus **it is very unlikely that any of the NP-complete problems would have a polynomial-time solution.**

However, currently **there is no proof to show either $P = NP$ or $P \neq NP$.** Namely, **the status of “ $P = NP$?” is still unknown**, and it is still the **biggest open problem** in Computer Science.

Elaboration: ScreenShots of Annotations

as a valid input, etc.) then w is not in the domain of x . It turns out that Turing machines have the same computing powers as computers (in terms of what functions can be computed on them, **not** in terms of the computing efficiencies).

Universal Turing Machine. A nice property is the **universality** of Turing machines, namely, we can have a **universal Turing machine** TM^* : given an input $\langle e(M), w \rangle$ where $e(M)$ is the **encoding** of some Turing machine M and w is a string, TM^* will **simulate** the actions of M on the input w . This essentially means that the universal Turing machine TM^* is a **general-purpose computer** and we can **program** it — the finite-length encoding $e(M)$ is some finite-length **program** (say in C++ or Java, etc.) and the simulation is to run the program on the input w .

Countable vs. Uncountable Sets

We say that a set A is **countable** if every element $a \in A$ can be mapped to a **distinct** number from the set \mathbb{N} of natural numbers $\{1, 2, 3, \dots\}$. For example, the set \mathbb{Z} of all integers is countable, since we can count them as follows: count/map $0, 1, -1, 2, -2, \dots$, in that order, as $1, 2, 3, 4, \dots$ (note: counting the integers in the order of $0, 1, 2, 3, \dots$ does **not** work, since this would only go along the positive integers and miss the negative ones). The set \mathbb{Q}^+ of positive rational numbers is countable, since such numbers can be represented as p/q where $p = 1, 2, 3, \dots$ and $q = 1, 2, 3, \dots$ (think of listing p/q in a 2-dimensional table where one dimension is indexed by p and the other by q). We start by counting the tuples (p, q) in the order of $(1, 1), (1, 2), (2, 1), (1, 3), (2, 2), (3, 1), \dots$. That is, with $p + q = 2$ there is one tuple $(1, 1)$; with $p + q = 3$ there are two tuples $(1, 2), (2, 1)$; with $p + q = 4$ there are three tuples $(1, 3), (2, 2), (3, 1)$, etc. (this is called **dovetailing**). In this way all tuples (p, q) are counted and no one is missed. The set \mathbb{Q}^- of all negative rational numbers can be counted symmetrically. Therefore, the set \mathbb{Q} of all rational numbers is countable, since we can first count 0, then the first number in \mathbb{Q}^+ , the first number in \mathbb{Q}^- , the second number in \mathbb{Q}^+ , the second number in \mathbb{Q}^- , and so on. For an infinite set, we say that it is **countably infinite** if it is countable; otherwise it is **uncountably infinite**.

Claim 1. The set of all Turing machines is countably infinite.

Proof: Each Turing machine M corresponds to its encoding $e(M)$, which is a finite-length string. Therefore we can sort these encodings/strings first by their length and then by their lexicographical

the set \mathbb{N} of natural numbers $\{1, 2, 3, \dots\}$. For example, the set \mathbb{Z} of all integers is countable, since we can count them as follows: count/map $0, 1, -1, 2, -2, \dots$, in that order, as $1, 2, 3, 4, \dots$ (note: counting the integers in the order of $0, 1, 2, 3, \dots$ does **not** work, since this would only go along the positive integers and miss the negative ones). The set \mathbb{Q}^+ of positive rational numbers is countable, since such numbers can be represented as p/q where $p = 1, 2, 3, \dots$ and $q = 1, 2, 3, \dots$ (think of listing p/q in a 2-dimensional table where one dimension is indexed by p and the other by q). We start by counting the tuples (p, q) in the order of $(1, 1), (1, 2), (2, 1), (1, 3), (2, 2), (3, 1), \dots$. That is, with $p + q = 2$ there is one tuple $(1, 1)$; with $p + q = 3$ there are two tuples $(1, 2), (2, 1)$; with $p + q = 4$ there are three tuples $(1, 3), (2, 2), (3, 1)$, etc. (this is called **dovetailing**). In this way all tuples (p, q) are counted and no one is missed. The set \mathbb{Q}^- of all negative rational numbers can be counted symmetrically. Therefore, the set \mathbb{Q} of all rational numbers is countable, since we can first count 0, then the first number in \mathbb{Q}^+ , the first number in \mathbb{Q}^- , the second number in \mathbb{Q}^+ , the second number in \mathbb{Q}^- , and so on. For an infinite set, we say that it is **countably infinite** if it is countable; otherwise it is **uncountably infinite**.

Claim 1. The set of all Turing machines is countably infinite.

Proof: Each Turing machine M corresponds to its encoding $e(M)$, which is a finite-length string. Therefore we can sort these encodings/strings first by their length and then by their lexicographical order within the group of the same length. We count these encodings/strings in this final sorted order as $1, 2, 3, \dots$. Clearly all strings are counted and thus the set is countably infinite. \square

Claim 2. The set of all real numbers in the range $[0, 1]$ is uncountably infinite.

Proof: The proof uses a proof technique called **diagonalization**. Consider representing such real numbers in binary, then each such $r \in [0, 1]$ is in the form $0.b_1 b_2 b_3 \dots$, where b_i is the i -th bit after

by appending bits of 0 if needed. Now consider the real number $r^* = 0.b_{11}b_{22}b_{33}b_{44}\dots b_{kk}\dots$, where $b_{ii}, i = 1, 2, \dots$ are the diagonal bits marked by the red lines in Fig. 2. Finally, we let $\bar{r}^* = 0.\bar{b}_{11}\bar{b}_{22}\bar{b}_{33}\bar{b}_{44}\dots\bar{b}_{kk}\dots$, where \bar{b}_{ii} means the complement of b_{ii} : if $b_{ii} = 0$ then $\bar{b}_{ii} = 1$; else ($b_{ii} = 1$) then $\bar{b}_{ii} = 0$. Clearly, \bar{r}^* differs from each r_1, r_2, \dots at the diagonal bit b_{ii} . However, \bar{r}^* is a real number in $[0, 1]$ and we must have $\bar{r}^* = r_k$ for some $k \in \mathbb{N}$, which is a contradiction since \bar{r}^* differs from r_k at b_{kk} . Therefore it is not possible to label all real numbers in $[0, 1]$ as r_1, r_2, \dots , meaning that the set in question is uncountable. \square

Note that the above proof shows that an uncountable set S has “more elements” than a countable set — when trying to count/label the elements of S by $1, 2, 3, \dots$ there is always some element of S “missing” that cannot be so counted/labeled.

Take-Away Message: There are more functions than Turing machines

For each real number $x \in [0, 1]$ we can define a **distinct function** $f_x : [0, 1] \rightarrow [0, 1]$ as follows: $f_x(0) = x$, and $f_x(t) = t$ for all $t \in [0, 1]$ and $t \neq 0$ (note that each f_x maps the same 0 to a **distinct** value x , and thus each f_x is **distinct**). Since by Claim 2 there are uncountably many real numbers $x \in [0, 1]$, there are also **uncountably many** functions f_x . However, by Claim 1, there are only countably many Turing machines. This means that there are **more functions than Turing machines**, and thus **there are functions that cannot be computed by Turing machines/computers**, i.e., there are problems that cannot be computed/solved/decided by a computer.

A famous undecidable/unsolvable problem is the **halting problem**: “Given a Turing machine encoding/computer program M and an input w , can M terminate on w ?” (For a general M and a general w , this is undecidable! This can be proved by diagonalization.)

2. NP-Completeness

Now we consider the problems that can be computed/solved/decided by Turing machines/computers, but the focus is on whether they can be done efficiently or not. In general, polynomial-time solutions are considered efficient, and solutions requiring more than polynomial time (e.g., exponential time) are considered inefficient.

(P stands for polynomial time). There is also a class NP , where NP stands for **non-deterministic polynomial time**. For this, we first need to revisit Turing machines. In the Turing machine discussed in the beginning of Sec. 1, each rule of action $((q, x) \rightarrow (q', y)$ or $(q, x) \rightarrow (q', m)$) describes a **deterministic action**, i.e., for each tuple (q, x) there is a **unique action**. Such Turing machine is called a **deterministic Turing machine**, and for each input w , the computation process is a **single path** with **polynomial length** for a deterministic Turing machine.

In a **non-deterministic Turing machine**, each rule of action is of the form $(q, x) \rightarrow S$, where S is a **finite set** of possible actions, i.e., $S = \{s_1, s_2, \dots, s_k\}$ for some fixed k and each s_i is either some (q', y) or some (q', m) . This means that for a given state q and the input symbol x read, it will **non-deterministically** choose some $s_i \in S$ as the action (note that we do not say **randomly** since there is **no probability involved**). For each input w , the computation process is a **tree** where each internal node has fan-out $|S|$ (the number of actions in S) if the corresponding rule of action is $(q, x) \rightarrow S$. The computation tree contains many paths; some may not even stop. Among these paths, if **there exists a path that stops in polynomial time** (i.e., **if there exists a path of polynomial length**), then we say that the input w takes **non-deterministic polynomial time**.

Definition of NP: The class NP is defined to be the set of problems that can be solved in non-deterministic polynomial time. Equivalently, it is the set of problems for which **a given solution can be verified in polynomial time**. For example, given a (large) integer N , factoring it is difficult, but given two factors, we can easily verify that $N = \dots$

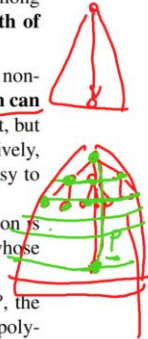
Turing machine.

In a **non-deterministic Turing machine**, each rule of action is of the form $(q, x) \rightarrow S$, where S is a **finite set** of possible actions, i.e., $S = \{s_1, s_2, \dots, s_k\}$ for some fixed k and each s_i is either some (q', y) or some (q', m) . This means that for a given state q and the input symbol x read, it will **non-deterministically** choose some $s_i \in S$ as the action (note that we do not say *randomly* since there is **no probability involved**). For each input w , the computation process is a **tree** where each internal node has fan-out $|S|$ (the number of actions in S) if the corresponding rule of action is $(q, x) \rightarrow S$. The computation tree contains many paths; some may not even stop. Among these paths, if **there exists a path that stops in polynomial time** (i.e., **if there exists a path of polynomial length**), then we say that the input w takes **non-deterministic polynomial time**.

Definition of NP: The class NP is defined to be the set of problems that can be solved in non-deterministic polynomial time. Equivalently, it is the set of problems for which **a given solution can be verified in polynomial time**. For example, given a (large) integer N , factoring it is difficult, but given two factors a and b as an answer, we can verify that $N = a \cdot b$ in polynomial time. Intuitively, P is a class of problems that are easy to compute, and NP is a class of problems that are easy to verify the solutions.

It is trivial to see that $P \subseteq NP$: from the deterministic Turing machine whose computation is a single path p of polynomial length, we can construct a non-deterministic Turing machine whose computation tree contains the path p .

All problems in NP have exponential time solutions. Recall that for any problem in NP , the corresponding computation tree of a non-deterministic Turing machine NTM has a path p of polynomial length ℓ . Then we can use a deterministic Turing machine TM to simulate NTM by exploring the computation tree in a **locked-step** fashion: in each round we advance the first path one more step, then the second path one more step, etc.; after the last path is advanced one more step, we repeat the process for the next round. This essentially **performs a BFS** in the computation tree; when we reach the leaf of path p then we find the solution and can stop. The total time is the **size of the computation tree cut at depth ℓ** ; if the maximum fan-out in the tree is k then such size is $O(k^\ell)$, i.e., the total time is exponential.



given two factors a and b as an answer, we can verify that $N = a \cdot b$ in polynomial time. Intuitively, P is a class of problems that are easy to compute, and NP is a class of problems that are easy to verify the solutions.

It is trivial to see that $P \subseteq NP$: from the deterministic Turing machine whose computation is a single path p of polynomial length, we can construct a non-deterministic Turing machine whose computation tree contains the path p .

All problems in NP have exponential time solutions. Recall that for any problem in NP , the corresponding computation tree of a non-deterministic Turing machine NTM has a path p of polynomial length ℓ . Then we can use a deterministic Turing machine TM to simulate NTM by exploring the computation tree in a **locked-step** fashion: in each round we advance the first path one more step, then the second path one more step, etc.; after the last path is advanced one more step, we repeat the process for the next round. This essentially performs a BFS in the computation tree; when we reach the leaf of path p then we find the solution and can stop. The total time is the **size of the computation tree cut at depth ℓ** ; if the maximum fan-out in the tree is k then such size is $O(k^\ell)$, i.e., the total time is exponential.

Optimization vs. Decision Problems. An optimization problem asks for a solution that *minimizes/maximizes* an objective function, while a decision problem asks for a yes/no question. Note that the **optimization** version of a problem is **at least as hard as** the corresponding decision version. For example, for the shortest-path problem, the optimization version asks for the length of a shortest path from s to t , while the decision version asks, for a given value k , whether there is a path from s to t with length at most k . Clearly, if we have the optimal length ℓ , we can answer yes/no by checking whether $\ell \leq k$.

To show that a problem is hard, it suffices to show that its **decision version is already hard** (then the optimization version is even harder). In the following, we only consider the decision problems.